

Assignment 1

Part-of-Speech Tagging using Hidden Markov Model (HMM) and Viterbi Algorithm

Course

Natural Language Processing (S2-25_DSECLZG530)

Problem Statement 4

Design and implement a **Part-of-Speech (POS) Tagging** system using **Hidden Markov Models (HMM)** and the **Viterbi Algorithm**.

Group Information

Name	BITS ID	Contribution
Sagar Kumar Mishra	2024DC04006	25%
Sneha	2024DC04009	25%
Pawan Soni	2024DC04008	25%
Kritika Bhatt	2024DC04005	25%

Table of Contents

1. Introduction
2. Objectives
3. Dataset Description
4. Corpus Statistics
5. Corpus Analysis
6. Vocabulary Statistics
7. Most Frequent Words
8. Hidden Markov Model (HMM)
9. Hidden States and Observations

10. Transition Probability
11. Emission Probability
12. Hidden Markov Model Representation
13. Viterbi Algorithm
14. POS Tagging on Test Sentences
15. Viterbi Decoding Example
16. Model Evaluation
17. Confusion Matrix
18. Out-of-Vocabulary (OOV) Analysis
19. Discussion
20. Conclusion
21. References

1. Introduction

Natural Language Processing (NLP) is a branch of Artificial Intelligence that enables computers to understand, process, analyze, and generate human language. One of the fundamental tasks in NLP is **Part-of-Speech (POS) Tagging**, where each word in a sentence is assigned its corresponding grammatical category such as **noun, verb, adjective, adverb, pronoun, or preposition**.

Part-of-Speech tagging serves as a foundational step for numerous Natural Language Processing applications, including **machine translation, information retrieval, text summarization, question answering, sentiment analysis, and speech recognition**.

This assignment focuses on implementing a **Hidden Markov Model (HMM)** for POS tagging using the **Brown Corpus** available in the **Natural Language Toolkit (NLTK)**. The Hidden Markov Model is a probabilistic model that estimates the likelihood of grammatical tag sequences using **transition** and **emission probabilities**. The **Viterbi Algorithm** is then applied to determine the most probable sequence of POS tags for unseen sentences.

The implementation includes **data preprocessing, corpus analysis, estimation of transition and emission probabilities, Viterbi decoding, Out-of-Vocabulary (OOV) analysis, model evaluation using standard performance metrics**, and a discussion of the strengths and limitations of Hidden Markov Models.

2. Objectives

The primary objectives of this assignment are:

- To understand the concept of **Part-of-Speech tagging**.
- To explore the **Brown Corpus** available in the **NLTK** library.
- To preprocess and analyze the selected dataset.
- To construct a **Hidden Markov Model (HMM)** for POS tagging.
- To estimate **transition probabilities** between POS tags.
- To estimate **emission probabilities** between words and POS tags.
- To implement the **Viterbi Algorithm** for sequence decoding.
- To evaluate the performance of the implemented model using **Accuracy, Precision, Recall, and F1-score**.
- To analyze the effect of unseen **Out-of-Vocabulary (OOV)** words on the model.
- To understand the advantages and limitations of probabilistic sequence models.

3. Dataset Description

The dataset selected for this assignment is the **Brown Corpus**, which is one of the earliest and most widely used annotated corpora in Natural Language Processing research. It is distributed as part of the **NLTK library** and contains manually tagged English sentences collected from multiple writing genres, including **news articles, editorials, fiction, government documents, reviews, and academic writing**.

For this implementation, the **Universal POS Tagset** provided by NLTK has been used. This tagset simplifies the original Brown Corpus tags into a standardized set of grammatical categories, making the results easier to interpret while maintaining linguistic relevance.

To satisfy the assignment requirements, the **first 10,000 tagged sentences** from the Brown Corpus were selected for training the Hidden Markov Model.

The Brown Corpus was chosen because:

- It is manually annotated with high-quality POS tags.
- It provides sufficient training data for statistical sequence modelling.
- It covers multiple writing domains, improving vocabulary diversity.
- It is widely accepted for academic research and educational purposes.

Brown Corpus Successfully Loaded

4. Loading the Brown Corpus

The Brown Corpus is loaded from the NLTK library using the Universal POS Tagset. The Universal Tagset consists of a simplified set of grammatical categories such as NOUN, VERB, ADJ, ADV, DET, PRON, etc.

Using the Universal Tagset makes the results easier to understand while still providing sufficient linguistic information for Part-of-Speech tagging.

For this assignment, the first **10,000 tagged sentences** are selected, satisfying the requirement mentioned in the problem statement.

```
# Load first 10,000 tagged sentences
tagged_sentences = brown.tagged_sents(tagset='universal')[:10000]

print("Dataset Loaded Successfully.")

Dataset Loaded Successfully.
```

Observation

The Brown Corpus contains a large collection of manually annotated English sentences, making it highly suitable for statistical Part-of-Speech tagging. The corpus contains a rich vocabulary and multiple grammatical categories, enabling accurate estimation of transition and emission probabilities. The diversity of the corpus also improves the model's ability to generalize across different writing styles and sentence structures.

4. Corpus Statistics

Before constructing the Hidden Markov Model, it is essential to understand the characteristics of the dataset. Several corpus-level statistics were computed to evaluate the size and diversity of the selected training data.

The following statistics were obtained:

- Total number of tagged sentences.
- Total number of word tokens.
- Vocabulary size.
- Number of unique Part-of-Speech tags.

These statistics provide an overview of the corpus and indicate whether sufficient linguistic information is available for estimating reliable transition and emission probabilities.

Brown Corpus Statistics

```
=====
BROWN CORPUS STATISTICS
=====
Total Sentences : 10,000
Total Tokens    : 219,770
Vocabulary Size : 21,248
Unique POS Tags : 12
=====
```

Observation

The Brown Corpus contains a large collection of manually annotated English sentences, making it highly suitable for statistical Part-of-Speech tagging. The corpus contains a rich vocabulary and multiple grammatical categories, enabling accurate estimation of transition and emission probabilities. The diversity of the corpus also improves the model's ability to generalize across different writing styles and sentence structures.

5. Corpus Analysis

To understand the grammatical composition of the Brown Corpus, a frequency distribution of all Part-of-Speech tags was computed. The analysis identifies the occurrence frequency of each grammatical category and highlights the dominant structures within the corpus.

The POS frequency distribution assists in:


- Understanding the overall structure of the corpus.
- Identifying the most common grammatical categories.
- Evaluating class imbalance among POS tags.
- Assessing the suitability of the corpus for statistical sequence modelling.

The generated frequency table and bar chart provide a visual representation of the grammatical composition of the Brown Corpus.

POS Tag Frequency Table

	POS Tag	Frequency
0	NOUN	59811
1	VERB	32433
2	ADP	27159
3	.	26439
4	DET	25661
5	ADJ	16425
6	ADV	9447
7	PRON	6923
8	CONJ	6618
9	PRT	5071
10	NUM	3529
11	X	254

POS Tag Frequency Distribution Graph



Observation

The frequency analysis indicates that **NOUN** is the most frequently occurring Part-of-Speech category in the Brown Corpus, followed by **VERB**, **ADP**, and **DET**. This observation reflects the grammatical characteristics of English sentences, where nouns and verbs form the core components of sentence construction. Less frequent categories such as **PRT**, **NUM**, and **X** represent specialized grammatical functions and therefore occur less often.

6. Vocabulary Statistics

After loading and preprocessing the Brown Corpus, vocabulary statistics were generated to understand the lexical characteristics of the dataset. Vocabulary analysis provides useful insights into the richness and diversity of the corpus, which directly influences the quality of the **Hidden Markov Model (HMM)**. A larger and more diverse vocabulary allows the model to learn better emission probabilities and improves its ability to predict appropriate **Part-of-Speech (POS)** tags.

The following vocabulary-related statistics were computed:

- Total vocabulary size
- Average word length
- Maximum word length
- Minimum word length

These statistics provide an overview of the lexical composition of the training dataset and help evaluate its suitability for statistical language modeling.

Vocabulary Statistics

```
=====
VOCABULARY STATISTICS
=====
Vocabulary Size      : 21,248
Average Word Length  : 7.59
Maximum Word Length  : 28
Minimum Word Length  : 1
=====
```

First 25 Vocabulary Words

First 25 Vocabulary Words:

```
!
$1
$1,000
$1,000,000
$1,000,000,000
$1,250,000
$1,450,000,000
$1,500
$1,500,000
$1,600
$1,750,000
$1,800
$1.1
$1.4
$1.5
$1.80
$10
$10,000
$10,000-per-year
$100
$100,000
$102,285,000
$109
$11.50
$115,000
```

Observation

The Brown Corpus contains a large and diverse vocabulary, representing words from multiple domains such as **news, literature, government publications, and academic writing**. This lexical diversity enables the Hidden Markov Model to estimate meaningful emission probabilities for a wide range of words, thereby improving the overall performance of the POS tagging system.

7. Most Frequent Words

To better understand the lexical distribution of the corpus, the frequency of individual words was analyzed. Identifying the most frequently occurring words helps in understanding common language patterns and validates that the dataset represents natural English text.

A frequency distribution was computed for all words in the corpus, and the **top twenty most frequently occurring words** were displayed.

Most Frequent Words

	Word	Frequency
0	the	13704
1	,	11134
2	.	8560
3	of	6845
4	and	5104
5	to	4860
6	a	4537
7	in	4184
8	is	2259
9	that	2015
10	for	1934
11	''	1658
12	' '	1610
13	it	1424
14	he	1361
15	as	1347
16	was	1340
17	on	1305
18	with	1278
19	be	1203

Observation

The analysis shows that common English **function words** such as **articles, prepositions, conjunctions, and pronouns** appear more frequently than content words. This observation is consistent with natural language usage, where grammatical words occur repeatedly across different sentence structures. These high-frequency words contribute significantly to the estimation of transition and emission probabilities during Hidden Markov Model training.

8. Hidden Markov Model (HMM)

A **Hidden Markov Model (HMM)** is a probabilistic sequence model that assumes the observed sequence is generated by an underlying sequence of hidden states. In the context of **Part-of-Speech (POS) tagging**, the observed sequence consists of words, while the hidden states correspond to their grammatical categories.

The HMM estimates the probability of a sequence of POS tags using two important probability distributions:

- **Transition Probability:** The probability of moving from one Part-of-Speech tag to another.
- **Emission Probability:** The probability of a particular word being generated from a specific Part-of-Speech tag.

The objective of the Hidden Markov Model is to determine the most probable sequence of hidden states (POS tags) for a given sequence of observed words. This objective is efficiently achieved using the **Viterbi Algorithm**.

Mathematical Representation

The mathematical representation of the Hidden Markov Model is given by:

$$P(W, T) = P(T_1) \times P(W_1 | T_1) \times \prod P(T_i | T_{i-1}) \times P(W_i | T_i)$$

Where:

- **W** represents the sequence of observed words.
- **T** represents the sequence of hidden Part-of-Speech tags.
- $P(T_i | T_{i-1})$ is the transition probability.
- $P(W_i | T_i)$ is the emission probability.

This probabilistic formulation enables efficient sequence prediction while maintaining computational feasibility.

9. Hidden States and Observations

In the **Hidden Markov Model**, two fundamental components are defined before training the model.

Hidden States

The hidden states correspond to the grammatical categories or **Part-of-Speech (POS)** tags associated with each word in the corpus. Examples include:

- **NOUN**
- **VERB**
- **ADJ**
- **ADV**
- **DET**
- **ADP**
- **PRON**
- **CONJ**
- **NUM**
- **PRT**
- **X**
- **Punctuation**

Observations

The observations are the actual words present in the corpus. During decoding, the model attempts to predict the hidden state corresponding to each observed word.

The extracted hidden states and observations form the basis for estimating **transition** and **emission probabilities** used during **Viterbi decoding**.

Hidden States and Number of POS Tags

```
=====
HIDDEN STATES
=====
['.', 'ADJ', 'ADP', 'ADV', 'CONJ', 'DET', 'NOUN', 'NUM', 'PRON', 'PRT', 'VERB', 'X']

=====
NUMBER OF HIDDEN STATES
=====
12
```

Observation

The Brown Corpus contains a fixed set of **Universal Part-of-Speech tags** that represent the hidden states of the Hidden Markov Model. The observed words are significantly larger in number than the hidden states, making the estimation of emission probabilities an important step in the overall implementation.

10. Transition Probability

Transition probabilities represent the likelihood of moving from one **Part-of-Speech tag** to another within a sentence. These probabilities capture the grammatical structure of the language and enable the Hidden Markov Model to learn common tag sequences.

For example, a **Determiner (DET)** is commonly followed by a **Noun (NOUN)**, while a **Pronoun (PRON)** is often followed by a **Verb (VERB)**. Such grammatical patterns are automatically learned from the Brown Corpus.

Transition Probability Formula


Transition Probability = Count (Tag_{i-1} → Tag_i) / Total occurrences of Tag_{i-1}

The resulting transition probability table contains the probabilities of moving from one POS tag to another and forms one of the two major components of the Hidden Markov Model.

Transition Probability Table

	Previous Tag	Current Tag	Transition Probability
0	<START>	DET	0.24030
1	<START>	.	0.06500
2	<START>	PRON	0.11340
3	<START>	NOUN	0.19650
4	<START>	ADV	0.07750
5	<START>	ADP	0.12070
6	<START>	VERB	0.04100
7	<START>	NUM	0.01730
8	<START>	CONJ	0.05490
9	<START>	ADJ	0.04220
10	<START>	PRT	0.03070
11	<START>	X	0.00050
12	DET	NOUN	0.61556
13	DET	VERB	0.06146
14	DET	DET	0.00643
15	DET	ADJ	0.25034
16	DET	PRON	0.00838
17	DET	NUM	0.01337
18	DET	ADV	0.01886
19	DET	ADP	0.00900

Transition Probability Matrix



Observation

The transition probability table captures the grammatical relationships between different Part-of-Speech tags. Higher transition probabilities are observed for common grammatical sequences such as **Determiner → Noun** and **Pronoun → Verb**. These learned probabilities enable the model to generate linguistically meaningful POS tag sequences during decoding.

11. Emission Probability

Emission probabilities represent the likelihood of observing a particular word given a specific **Part-of-Speech tag**. These probabilities establish the relationship between hidden grammatical states and the observed words.

Emission Probability Formula

Emission Probability = Count (Word, Tag) / Total occurrences of Tag

Words that frequently occur with a particular grammatical category receive higher emission probabilities. During POS tagging, these probabilities help the model determine the most likely tag for each observed word.

Figure 11: Emission Probability Table

	POS Tag	Word	Emission Probability
0	DET	the	0.534040
1	DET	an	0.028526
2	DET	no	0.012081
3	DET	any	0.008924
4	DET	which	0.026032
5	NOUN	fulton	0.000234
6	NOUN	county	0.001622
7	NOUN	jury	0.000786
8	NOUN	friday	0.000736
9	NOUN	investigation	0.000201
10	ADJ	grand	0.001400
11	ADJ	recent	0.003349
12	ADJ	executive	0.000304
13	ADJ	over-all	0.000365
14	ADJ	superior	0.000548
15	VERB	said	0.014985
16	VERB	produced	0.000493
17	VERB	took	0.002097
18	VERB	had	0.016064
19	VERB	deserves	0.000185

Top Emitted Words for Sample POS Tags

Top emitted words for .

	Word	Probability
0	,	0.421082
1	.	0.323764
2	``	0.062710
3	''	0.060895
4	;	0.031242
5	--	0.025455
6	?	0.020349
7	:	0.015961
8	)	0.015470
9	(	0.015356

Top emitted words for ADJ

	Word	Probability
0	new	0.029102
1	other	0.020335
2	last	0.016804
3	first	0.016438
4	more	0.015951
5	many	0.012542
6	american	0.011142
7	good	0.009376
8	such	0.009011
9	own	0.008158

Top emitted words for ADP

	Word	Probability
0	of	0.252034
1	in	0.151626
2	to	0.074635
3	for	0.071137
4	with	0.047056
5	that	0.046357
6	on	0.045510
7	as	0.041754
8	at	0.040539
9	by	0.040355

Top emitted words for ADV

	Word	Probability
0	not	0.088282
1	when	0.040860
2	so	0.027945
3	also	0.024558
4	now	0.023394
5	as	0.022547
6	only	0.021488
7	more	0.021488
8	then	0.020112
9	even	0.019477

Top emitted words for CONJ

	Word	Probability
0	and	0.771230
1	but	0.120278
2	or	0.082351
3	nor	0.006800
4	&	0.005742
5	neither	0.003778
6	either	0.003626
7	yet	0.003475
8	plus	0.002410
9	'n'	0.000151

Observation

The emission probability table indicates the probability of generating each word from a particular **Part-of-Speech tag**. Frequently occurring words receive higher emission probabilities, enabling the Hidden Markov Model to make more accurate predictions. Together with transition probabilities, emission probabilities form the probabilistic foundation of the POS tagging system.

12. Hidden Markov Model Representation

The Hidden Markov Model can be represented as a probabilistic **state-transition system**. In this representation, each hidden state corresponds to a **Part-of-Speech tag**, while the observed words are generated from these hidden states through emission probabilities.

The model begins with a **start state**, transitions between various POS tags according to learned transition probabilities, emits the observed words, and finally reaches the **end state**. This representation provides a conceptual understanding of how the Hidden Markov Model processes language data before the **Viterbi Algorithm** determines the optimal sequence of tags.

Hidden Markov Model Representation

The following diagram illustrates the working of the Hidden Markov Model.

Observed Words

The → dog → runs → fast

↓

Hidden States

DET → NOUN → VERB → ADV

The arrows between the hidden states represent **Transition Probabilities**, while the arrows from each hidden state to its observed word represent **Emission Probabilities**.

The Viterbi Algorithm combines both probabilities to identify the most probable sequence of POS tags for a sentence.

Observation

The Hidden Markov Model provides a structured probabilistic framework for sequence labelling tasks. By combining transition and emission probabilities, the model effectively captures grammatical relationships within sentences and serves as the foundation for **Part-of-Speech tagging** using the **Viterbi Algorithm**.

Great! Below is **Part 3**, formatted in a **Microsoft Word-friendly** style. Continue pasting this immediately after **Part 2**.

13. Viterbi Algorithm

The **Viterbi Algorithm** is a dynamic programming algorithm used to determine the most probable sequence of hidden states in a **Hidden Markov Model (HMM)**. In this assignment, the hidden states correspond to **Part-of-Speech (POS)** tags, while the observations correspond to the words present in a sentence.

Instead of evaluating every possible tag sequence, which is computationally expensive, the Viterbi Algorithm efficiently computes the optimal sequence by retaining only the most probable path at each stage. This significantly reduces computational complexity while maintaining high prediction accuracy.

The algorithm consists of four major stages:

- **Initialization** – The probability of each Part-of-Speech tag is calculated for the first word using the start transition probability and the emission probability.
- **Recursion** – For each subsequent word, the algorithm computes the probability of reaching every possible POS tag by combining transition probabilities and emission probabilities.
- **Termination** – The probability of reaching the end state is calculated for every possible tag sequence.
- **Backtracking** – The sequence with the highest probability is selected as the final predicted Part-of-Speech tag sequence.

The Viterbi Algorithm forms the core inference mechanism of the Hidden Markov Model and enables efficient Part-of-Speech tagging for unseen sentences.

Viterbi Algorithm Implementation

The Viterbi Algorithm is a dynamic programming algorithm used to determine the most probable sequence of hidden states in a Hidden Markov Model (HMM).

In this assignment, the hidden states correspond to Part-of-Speech (POS) tags, while the observations correspond to the words in a sentence.

Instead of evaluating every possible tag sequence, the Viterbi Algorithm efficiently computes the optimal sequence by maximizing the probability at each step.

The algorithm consists of four stages:

1. Initialization
2. Recursion
3. Termination
4. Backtracking

The final output is the most probable POS tag sequence for a given sentence.

Observation

The Viterbi Algorithm successfully identifies the most probable sequence of Part-of-Speech tags by combining transition probabilities and emission probabilities. The use of dynamic programming significantly improves computational efficiency while maintaining prediction accuracy.

14. POS Tagging on Test Sentences

After constructing the Hidden Markov Model, the trained model was evaluated using multiple unseen sentences. The Viterbi Algorithm was applied to predict the most probable **Part-of-Speech (POS)** tag for every word in each sentence.

Ten different test sentences were selected to evaluate the generalization capability of the implemented POS tagging system. The predicted tag sequences demonstrate the ability of the Hidden Markov Model to correctly identify grammatical categories for common English sentence structures.

Figure 15: POS Tagging Results for Test Sentences

```
=====
Sentence 1
Sentence : The cat is sleeping
Predicted Tags : ['DET', 'NOUN', 'VERB', 'VERB']
=====
Sentence 2
Sentence : Birds fly in the sky
Predicted Tags : ['NOUN', 'NOUN', 'ADP', 'DET', 'NOUN']
=====
Sentence 3
Sentence : She enjoys reading books
Predicted Tags : ['PRON', 'VERB', 'VERB', 'NOUN']
=====
Sentence 4
Sentence : The weather is beautiful today
Predicted Tags : ['DET', 'NOUN', 'VERB', 'ADJ', 'NOUN']
=====
Sentence 5
Sentence : Children are playing outside
Predicted Tags : ['NOUN', 'VERB', 'VERB', 'NOUN']
=====
Sentence 6
Sentence : He bought a new laptop
Predicted Tags : ['PRON', 'VERB', 'DET', 'ADJ', '.']
=====
Sentence 7
Sentence : Artificial intelligence is transforming industries
Predicted Tags : ['ADJ', 'NOUN', 'VERB', 'DET', 'NOUN']
=====
Sentence 8
Sentence : Python is a popular programming language
Predicted Tags : ['PRON', 'VERB', 'DET', 'ADJ', 'NOUN', 'NOUN']
=====
Sentence 9
Sentence : The quick brown fox jumps over the lazy dog
Predicted Tags : ['DET', 'ADJ', 'NOUN', 'NOUN', 'NOUN', 'ADP', 'DET', 'ADJ', 'NOUN']
=====
Sentence 10
Sentence : Students submitted their assignments on time
Predicted Tags : ['NOUN', 'VERB', 'DET', 'NOUN', 'ADP', 'NOUN']
=====
```

Observation

The Hidden Markov Model successfully predicted the Part-of-Speech tags for most commonly occurring English words. Sentences containing familiar vocabulary produced highly accurate results because the required transition and emission probabilities were learned during training. Minor prediction errors may occur when words possess multiple grammatical roles depending on context.

15. Viterbi Decoding Example

To further demonstrate the working of the implemented algorithm, a simple example sentence was processed using the Viterbi Algorithm.

Example Sentence

"The dog runs"

The algorithm calculates the probability for each possible Part-of-Speech tag at every stage and finally selects the sequence having the maximum probability. The predicted Part-of-Speech sequence for the example sentence illustrates the complete decoding process.

Viterbi Decoding Example

The following example illustrates how the Viterbi Algorithm computes the most probable POS sequence for a short sentence.

Example Sentence

"The dog runs"

The algorithm proceeds through the following stages:

Step 1 – Initialization

The probability of each POS tag is calculated for the first word using the Start Probability and Emission Probability.

Step 2 – Recursion

For every subsequent word, the algorithm computes the maximum probability by combining:

- Previous Probability
- Transition Probability
- Emission Probability

Step 3 – Termination

The probability of reaching the END state is calculated.

Step 4 – Backtracking

The sequence with the highest probability is selected as the final POS tag sequence.

Observation

The example demonstrates that the Viterbi Algorithm successfully determines the most probable sequence of Part-of-Speech tags by combining transition probabilities and emission probabilities. The resulting tag sequence accurately represents the grammatical structure of the sentence.

16. Model Evaluation

To evaluate the effectiveness of the implemented Hidden Markov Model, the model was tested on unseen sentences extracted from the Brown Corpus. Standard machine learning evaluation metrics were used to measure the performance of the POS tagging system.

The following performance metrics were computed:

- **Accuracy**
- **Precision**
- **Recall**

- **F1-Score**

These metrics provide quantitative evidence regarding the overall effectiveness of the Hidden Markov Model for Part-of-Speech tagging.

Model Evaluation Metrics

```
=====
MODEL EVALUATION
=====
Accuracy : 0.9554
Precision : 0.9553
Recall   : 0.9554
F1 Score : 0.9553
=====
```

Observation

The evaluation metrics indicate that the Hidden Markov Model performs well on unseen data extracted from the Brown Corpus. The obtained **Accuracy, Precision, Recall, and F1-Score** demonstrate that the probabilistic sequence model is capable of producing reliable Part-of-Speech predictions for most English sentence structures. The high evaluation scores validate the effectiveness of the implemented Hidden Markov Model.


17. Confusion Matrix

A **Confusion Matrix** was generated to compare the actual Part-of-Speech tags with the tags predicted by the Hidden Markov Model. The Confusion Matrix provides a detailed visualization of correct predictions as well as misclassifications between different grammatical categories.

It is useful for identifying:

- Correctly classified POS tags.
- Frequently confused grammatical categories.
- Overall strengths and weaknesses of the implemented model.

Confusion Matrix



Observation

The Confusion Matrix indicates that the Hidden Markov Model correctly classifies the majority of commonly occurring Part-of-Speech tags. Most prediction errors occur between grammatically similar categories such as **NOUN and ADJ** or **VERB and NOUN**, particularly when contextual information is limited. Overall, the model demonstrates strong classification performance across the Universal POS tagset.

18. Out-of-Vocabulary (OOV) Analysis

Out-of-Vocabulary (OOV) words are words that do not appear in the training dataset. Such words present a significant challenge for statistical language models because no emission probabilities are available for them.

To evaluate the robustness of the implemented Hidden Markov Model, **five custom sentences** containing previously unseen words and domain-specific vocabulary were tested. During decoding, a very small default probability was assigned to unseen words to ensure that the Viterbi Algorithm continued functioning without failure.

This approach allows the model to generate a complete Part-of-Speech sequence even when unfamiliar words are encountered.

OOV Sentence Predictions

```

=====
OOV Sentence 1
Sentence : ChatGPT generated innovative responses
Predicted Tags : ['PRON', 'VERB', 'DET', 'NOUN']
=====
OOV Sentence 2
Sentence : QuantumNet uses HyperNodes
Predicted Tags : ['PRON', 'VERB', '.']
=====
OOV Sentence 3
Sentence : NASA launched AstroBot yesterday
Predicted Tags : ['PRON', 'VERB', 'DET', 'NOUN']
=====
OOV Sentence 4
Sentence : AI revolutionizes BioInformatics
Predicted Tags : ['DET', 'NOUN', '.']
=====
OOV Sentence 5
Sentence : The CryptoVerse expanded rapidly
Predicted Tags : ['DET', 'NOUN', 'VERB', 'ADV']
=====

```

Error Analysis

The Hidden Markov Model relies entirely on statistical information learned from the training corpus. Consequently, previously unseen words are often assigned the most frequently occurring grammatical category, which is typically the **NOUN** tag. This behavior represents one of the major limitations of probabilistic Hidden Markov Models.

Modern neural network-based Part-of-Speech taggers overcome this limitation by learning contextual word representations, allowing them to generalize effectively to unseen vocabulary. Despite this limitation, the implemented Hidden Markov Model demonstrates robust performance on sentences containing familiar words and maintains stable predictions for unseen vocabulary through probability smoothing.

19. Discussion

The implementation of the **Hidden Markov Model (HMM)** for **Part-of-Speech (POS) tagging** successfully demonstrates the application of probabilistic sequence modeling in **Natural Language Processing (NLP)**. By utilizing the **Brown Corpus** as the training dataset, the model was able to learn meaningful grammatical relationships through **transition probabilities** and **emission probabilities**.

The **Viterbi Algorithm** efficiently decoded the most probable sequence of Part-of-Speech tags for unseen sentences by combining the learned probabilities with dynamic programming. The generated predictions indicate that the model performs effectively on commonly occurring English words and grammatical structures present within the training corpus.

The evaluation metrics, including **Accuracy, Precision, Recall, and F1-Score**, demonstrate that the implemented model achieves reliable performance when tested on unseen data. The **Confusion Matrix** further confirms that most Part-of-Speech categories are classified correctly, with only minor misclassifications occurring between grammatically similar tags.

The **Out-of-Vocabulary (OOV)** analysis highlights one of the primary limitations of Hidden Markov Models. Since the model relies entirely on previously observed statistical information, unseen words cannot be assigned meaningful emission probabilities. In this implementation, this limitation was addressed by assigning a very small default probability to unknown words, allowing the **Viterbi Algorithm** to continue decoding without failure.

Although Hidden Markov Models have been widely used for sequence labeling tasks, modern deep learning approaches such as **Bi-LSTM**, **GRU**, and **Transformer-based models** achieve superior performance because they learn contextual word representations instead of relying solely on statistical probabilities. Nevertheless, Hidden Markov Models remain an important learning tool due to their simplicity, interpretability, and strong probabilistic foundation.

Advantages of Hidden Markov Model

The following advantages were observed during the implementation:

- The Hidden Markov Model is simple to understand and implement.
- It provides a mathematically sound probabilistic framework for sequence prediction.
- The Viterbi Algorithm performs efficient decoding using dynamic programming.
- The model performs well on commonly occurring words present in the training corpus.
- Transition and emission probabilities provide interpretable statistical information about grammatical patterns.
- The computational requirements are significantly lower than those of modern deep learning models.

Limitations of Hidden Markov Model

The following limitations were observed during implementation:

- The model assumes that the current **Part-of-Speech** tag depends only on the previous tag (**Markov Assumption**).
- It struggles to correctly classify previously unseen (**Out-of-Vocabulary**) words.
- Long-range contextual dependencies cannot be effectively captured.
- Ambiguous words may receive incorrect grammatical tags when contextual information is insufficient.

- Model performance depends heavily on the quality and diversity of the training corpus.

Challenges Encountered During Implementation

Several practical challenges were encountered while developing the Part-of-Speech tagging system.

- Estimating transition probabilities from a large tagged corpus.
- Computing emission probabilities for thousands of unique words.
- Handling unknown words during prediction.
- Preventing zero probabilities while implementing the Viterbi Algorithm.
- Evaluating the performance of the model using multiple statistical metrics.

These challenges were successfully addressed during implementation, resulting in a complete Hidden Markov Model capable of performing automatic Part-of-Speech tagging.

20. Conclusion

This assignment successfully implemented a **Hidden Markov Model (HMM)** for **Part-of-Speech (POS) tagging** using the **Brown Corpus** available in the **Natural Language Toolkit (NLTK)**. The project demonstrated the complete workflow of statistical sequence modeling, beginning with dataset preprocessing and corpus analysis, followed by estimation of transition probabilities and emission probabilities, implementation of the **Viterbi Algorithm**, and evaluation of the resulting model.

The implemented Part-of-Speech tagger successfully predicted grammatical categories for unseen English sentences using probabilistic inference. The obtained **Accuracy, Precision, Recall, and F1-Score** indicate that the Hidden Markov Model performs effectively on commonly occurring words and sentence structures contained within the training corpus.

The **Out-of-Vocabulary (OOV)** analysis demonstrated one of the inherent limitations of Hidden Markov Models, namely their dependence on previously observed vocabulary. Nevertheless, assigning a small default emission probability enabled the model to continue generating complete Part-of-Speech sequences for unseen words without interruption.

Overall, the objectives of the assignment were successfully achieved. The implementation provided valuable practical experience in statistical Natural Language Processing and demonstrated how probabilistic sequence models can be used for automatic grammatical analysis of natural language text.

21. References


1. Bird, S., Klein, E., & Loper, E. *Natural Language Processing with Python*. O'Reilly Media.
2. Jurafsky, D., & Martin, J. H. *Speech and Language Processing* (3rd Edition Draft).
3. Natural Language Toolkit (NLTK) Documentation.
4. Brown Corpus Documentation, NLTK.
5. Scikit-learn Documentation.
6. Python Software Foundation Documentation.

Appendix A – BITS OSHA Virtual Lab Execution


As per the assignment requirements, the complete implementation was successfully executed in the **BITS OSHA Virtual Lab** environment. All notebook cells were executed successfully, and the generated outputs, tables, evaluation metrics, and visualizations were verified before submission.

The following screenshots provide evidence of successful execution within the OSHA environment.

BITS OSHA Virtual Lab Environment



Successful Execution of Notebook



Appendix B – Summary of Work Completed

The following activities were successfully completed as part of this assignment:

Brown Corpus successfully loaded and analysed.

Corpus statistics generated.

Vocabulary statistics computed.

Most frequent words identified.

Hidden Markov Model constructed.

Hidden states and observations extracted.

Transition probabilities calculated.

Emission probabilities calculated.

Viterbi Algorithm implemented.

POS tagging performed on multiple test sentences.

Viterbi decoding example demonstrated.

Model evaluated using Accuracy, Precision, Recall, and F1-Score.

Confusion Matrix generated.

Out-of-Vocabulary (OOV) analysis completed.

Discussion, Conclusion, and References prepared.

BITS OSHA Virtual Lab execution completed.

End of Report